



# Writing Custom Tracing Functions

Tom Offermann  
October 1, 2016  
PyDX



# Where to Find Me

- ▶ **Work:** New Relic
- ▶ **Github:** toffer
- ▶ **Twitter:** toffermann
- ▶ **Email:** tom@offermand.us
- ▶ **Web:** www.offermand.us



# What is Tracing?

A "running narrative" of the line-by-line execution of your code.



# Why Tracing?

- ▶ Understand your program
- ▶ Powerful, but underutilized tool
- ▶ Better understand Python



# Agenda

1. Simple Tracing
2. How Does Tracing Work?
3. Custom Tracing
4. Who Else Uses Custom Tracing?
5. Full-Fledged Example



# Agenda

- 1. Simple Tracing**
2. How Does Tracing Work?
3. Custom Tracing
4. Who Else Uses Custom Tracing?
5. Full-Fledged Example

# Example: parent\_child.py

```
def child():
    for i in range(3):
        print(' ' * 8, 'in child loop')

def parent():
    for i in range(2):
        print(' ' * 4, 'in parent loop')
        child()

if __name__ == '__main__':
    print('main start')
    parent()
    print('main end')
```



# Output: parent\_child.py

```
$ python parent_child.py
```

```
main start
    in parent loop
        in child loop
        in child loop
        in child loop
    in parent loop
        in child loop
        in child loop
        in child loop
main end
```





# How to Trace

```
$ python -m trace --trace parent_child.py
```

# Trace Output


```
--- module: parent_child, funcname: <module>
parent_child.py(3): def child():
parent_child.py(7): def parent():
parent_child.py(12): if __name__ == '__main__':
parent_child.py(13):     print('main start')
main start
parent_child.py(14):     parent()
    --- module: parent_child, funcname: parent
parent_child.py(8):     for i in range(2):
parent_child.py(9):         print(' ' * 4, 'in parent loop')
    in parent loop
...
```

# Trace Output

```
--- module: parent_child, funcname: <module>
parent_child.py(3): def child():
parent_child.py(7): def parent():
parent_child.py(12): if __name__ == '__main__':
parent_child.py(13):     print('main start')
main start
parent_child.py(14):     parent()
    --- module: parent_child, funcname: parent
parent_child.py(8):     for i in range(2):
parent_child.py(9):         print(' ' * 4, 'in parent loop')
    in parent loop
...
```

# Trace Output

```
--- module: parent_child, funcname: <module>
parent_child.py(3): def child():
parent_child.py(7): def parent():
parent_child.py(12): if __name__ == '__main__':
parent_child.py(13):     print('main start')
main start
parent_child.py(14):     parent()
    --- module: parent_child, funcname: parent
parent_child.py(8):     for i in range(2):
parent_child.py(9):         print(' ' * 4, 'in parent loop')
    in parent loop
...
```



# Example: teapot.py

```
import requests
```

```
def teapot():  
    url = 'http://httpbin.org/status/418'  
    resp = requests.get(url)  
    print(resp.status_code, resp.reason)  
    print(resp.text)
```

```
teapot()
```

# Output: teapot.py

```
$ python teapot.py
```

```
418 I'M A TEAPOT
```

```
-=[ teapot ]=-
```

```
      _ . . . _  
      ,         \  
      .   _ _   .  
      |   " \ ^ \ "  
      \ . " _ _ " \ . _ ,  
      \ _ ; " _ _ _ " \ | //  
      |         ; /  
      \ _ _ _ /  
      \ " " " "  
      \ " " " "
```



# Trace teapot.py

```
$ python -m trace --trace teapot.py
```

# Trace Output (Start)

```
--- module: origin, funcname: <module>
origin.py(3): import requests
--- module: _bootstrap, funcname: _find_and_load
<frozen importlib._bootstrap>(968): --- module:
_bootstrap, funcname: __init__
<frozen importlib._bootstrap>(160): <frozen
importlib._bootstrap>(161): --- module: _bootstrap,
funcname: __enter__
<frozen importlib._bootstrap>(164): <frozen
importlib._bootstrap>(165): --- module: _bootstrap,
funcname: _get_module_lock
...
```



# Trace Output (End)

```
models.py(801):          return content
    -=[ teapot ]=-
```

```

      . . . .
    , - . . . - \
    .   - -   .
  | .  " \ ^ \ " .
  \_ ; " --- " \ | //
    |              ; /
    \_            _/
      \ " " " " \
```


```
--- modulename: response, funcname: closed
response.py(408):          if self._fp is None:
response.py(410):          elif hasattr(self._fp, 'closed'):
response.py(411):              return self._fp.closed
--- modulename: trace, funcname: _unsettrace
trace.py(77):              sys.settrace(None)
```



# Standard Library

```
$ python -m trace --trace teapot.py |  
>      grep client.py | head -5
```

```
client.py(69): """  
client.py(71): import email.parser  
<frozen ...: client.py(72): import email.message  
client.py(73): import http  
client.py(74): import io
```



# Requests

```
$ python -m trace --trace teapot.py |  
>      grep api.py | head -5
```

```
api.py(11): """
```

```
api.py(13): from . import sessions
```

```
<frozen ...: api.py(16): def request(method, ...)
```

```
api.py(59): def get(url, params=None, **kwargs):
```

```
api.py(73): def options(url, **kwargs):
```



# Lots of Trace Output

```
$ python -m trace --trace teapot.py | wc -l
```

```
250230
```

# Lots of Trace Output

```
$ python -m trace --trace teapot.py |  
>     grep -v '----' |                # Ignore function entry  
>     cut -f 1 -d '(' |                # Grab module name  
>     sort | uniq -c |                # Count by module  
>     sort -rn | head -n 5           # Show top 5
```

```
120836 sre_parse.py  
36957  sre_compile.py  
11800  enum.py  
4325   ipaddress.py  
3245   entities.py
```

# Example: teapot\_trace.py

```
import requests
from trace import Trace

def teapot():
    url = 'http://httpbin.org/status/418'
    resp = requests.get(url)
    print(resp.status_code, resp.reason)
    print(resp.text)


Trace().runfunc(teapot)
```



# More Focused Results

```
$ python teapot_trace.py | wc -l
```

```
13829
```



# Summary: Simple Tracing

- ▶ Built in, so always available
- ▶ Can run it from command line
- ▶ Extremely verbose!
- ▶ Use **Trace()** for focused results





# Agenda

1. Simple Tracing
- 2. How Does Tracing Work?**
3. Custom Tracing
4. Who Else Uses Custom Tracing?
5. Full-Fledged Example



# How Does Tracing Work?

1. Create a tracing function
2. Register the tracing function
3. Python will run tracer before executing each line of code.



# Create a Tracing Function



# Trace Line

```
class Trace:
```

```
...
```

```
def localtrace_trace(self, frame, why, arg):
```

```
    if why == "line":
```

```
        filename = frame.f_code.co_filename
```

```
        bname = os.path.basename(filename)
```

```
        lineno = frame.f_lineno
```

```
        line = linecache.getline(filename, lineno)
```

```
        print("%s(%d): %s" % (bname, lineno, line))
```

```
    return self.localtrace_trace
```

# Trace Line: Arguments

```
class Trace:
```

```
...
```

```
def localtrace_trace(self, frame, why, arg):
```

```
    if why == "line":
```

```
        filename = frame.f_code.co_filename
```

```
        bname = os.path.basename(filename)
```

```
        lineno = frame.f_lineno
```

```
        line = linecache.getline(filename, lineno)
```

```
        print("%s(%d): %s" % (bname, lineno, line))
```

```
    return self.localtrace_trace
```

# Trace Line: Return value

```
class Trace:
```

```
...
```

```
def localtrace_trace(self, frame, why, arg):
```

```
    if why == "line":
```

```
        filename = frame.f_code.co_filename
```

```
        bname = os.path.basename(filename)
```

```
        lineno = frame.f_lineno
```

```
        line = linecache.getline(filename, lineno)
```

```
        print("%s(%d): %s" % (bname, lineno, line))
```

```
    return self.localtrace_trace
```

# Argument #1: Frame

```
class Trace:
```

```
...
```

```
def localtrace_trace(self, frame, why, arg):  
    if why == "line":  
        filename = frame.f_code.co_filename  
        bname = os.path.basename(filename)  
        lineno = frame.f_lineno  
        line = linecache.getline(filename, lineno)  
        print("%s(%d): %s" % (bname, lineno, line))  
    return self.localtrace_trace
```



# **What is a Frame?**

## **(Detour ahead!)**



# How Python executes code

```
def parent():  
    for i in range(3):  
        child()  
    return 'parent done'
```

```
b'x\x1b\x00t\x00\x00d\x01\x00\x83\x01\x00D]\r  
\x00}\x00\x00t\x01\x00\x83\x00\x00\x01q\r\x00Wd  
\x02\x00S'
```



# Code Object

## Code

parent

`b'x\x1b...\x00S'`



# Frame Object

**Frame**

**Code**

parent

b'x\x1b...\x00S'



# Stack

## Frame

### Code

parent

`b'x\x1b...\x00S'`

## Frame

### Code

child

`b'x\x1e...\x00S'`



# Summary: Frames

- ▶ Source code compiled to bytes
- ▶ Python interpreter runs byte-compiled code
- ▶ Byte-compiled code is in "code object"
- ▶ Code object is in "execution frame"
- ▶ New frame created for every function call.
- ▶ Frames exist on "Stack"

# Stack Visualization

Python 3.3

```
→ 1 def child():
  2     for i in range(3):
  3         print(' ' * 8, 'in child loop')
  4
  5 def parent():
  6     for i in range(2):
  7         print(' ' * 4, 'in parent loop')
  8         child()
  9
10 if __name__ == '__main__':
11     print('main start')
12     parent()
13     print('main end')
```

[Edit code](#) | [Live programming](#)

Print output (drag lower right corner to resize)



Frames

Objects

# pythontutor.com #1

Print output (drag lower right corner to resize)

```
main start
```

Frames

Objects

Global frame

child

parent

function  
child()

function  
parent()



# pythontutor.com #2

Print output (drag lower right corner to resize)

```
main start  
  in parent loop
```

Frames

Objects

Global frame

child

parent

function  
child()

function  
parent()

parent

i | 0



# pythontutor.com #3

Print output (drag lower right corner to resize)

```
main start
  in parent loop
    in child loop
    in child loop
    in child loop
```

Frames

Objects

Global frame

child

parent

function  
child()

function  
parent()

parent

i | 0

child

i | 2

Return  
value

None

# pythontutor.com #4

Print output (drag lower right corner to resize)

```
main start
  in parent loop
    in child loop
    in child loop
    in child loop
```

Frames

Objects

Global frame

child

parent

function  
child()

function  
parent()

parent

i | 0

# pythontutor.com #5

Print output (drag lower right corner to resize)

```
main start
  in parent loop
    in child loop
    in child loop
    in child loop
  in parent loop
    in child loop
    in child loop
    in child loop
```

Frames

Objects

Global frame

child

parent

function  
child()

function  
parent()

parent

i | 1

Return  
value

None


# pythontutor.com #6

Print output (drag lower right corner to resize)

```
main start
  in parent loop
    in child loop
    in child loop
    in child loop
  in parent loop
    in child loop
    in child loop
    in child loop
main end
```

Frames

Objects





**Frames are objects**

# Frame Attributes

```
>>> help(sys._getframe())
```

```
class frame(object)
```

```
...
```

f_back	# this frame's caller
f_builtins	# builtins seen by this frame
f_code	# code object being executed
f_globals	# globals seen by this frame
f_lasti	# index of last instruction
f_lineno	# current line number in source
f_locals	# locals seen by this frame
f_trace	# tracing function for this frame



**Code objects are objects, too!**

# Code Object Attributes

```
>>> help(sys._getframe().f_code)
```

```
class code(object)
```

```
...
```

co_argcount	# Number of args
co_code	# String of bytecode
co_filename	# Name of file
co_firstlineno	# First source line number
co_name	# Name of code object
co_names	# Tuple of local var names
co_nlocals	# Number of local vars
co_varnames	# Names of args and local vars





**Detour Finished!**

# Arguments #2 & 3: Why and Arg

```
class Trace:
```

```
...
```

```
def localtrace_trace(self, frame, why, arg):
```

```
    if why == "line":
```

```
        filename = frame.f_code.co_filename
```

```
        bname = os.path.basename(filename)
```

```
        lineno = frame.f_lineno
```

```
        line = linecache.getline(filename, lineno)
```

```
        print("%s(%d): %s" % (bname, lineno, line))
```

```
    return self.localtrace_trace
```



# Why and Arg

Why (or Event)	Argument Value
call	None
line	None
return	Value being returned
exception	The (exception, val, tb) tuple

Source:

Doug Hellmann's Python 3 Module of the Week Series

- <https://pymotw.com/3/sys/tracing.html>



# Using Frame, Event and Arg

# Trace Line

```
def localtrace_trace(self, frame, why, arg):
    if why == "line":
        filename = frame.f_code.co_filename
        bname = os.path.basename(filename)
        lineno = frame.f_lineno
        line = linecache.getline(filename, lineno)

        print("%s(%d): %s" % (bname, lineno, line),
              end=' ')

    return self.localtrace
```

# Check Event

```
def localtrace_trace(self, frame, why, arg):  
    if why == "line":  
        filename = frame.f_code.co_filename  
        bname = os.path.basename(filename)  
        lineno = frame.f_lineno  
        line = linecache.getline(filename, lineno)  
  
        print("%s(%d): %s" % (bname, lineno, line),  
              end=' ')  
  
    return self.localtrace
```

# Access frame attribute

```
def localtrace_trace(self, frame, why, arg):  
    if why == "line":  
        filename = frame.f_code.co_filename  
        bname = os.path.basename(filename)  
        lineno = frame.f_lineno  
        line = linecache.getline(filename, lineno)  
  
        print("%s(%d): %s" % (bname, lineno, line),  
              end=' ')  
  
    return self.localtrace
```

# Access code attribute

```
def localtrace_trace(self, frame, why, arg):  
    if why == "line":  
        filename = frame.f_code.co_filename  
        bname = os.path.basename(filename)  
        lineno = frame.f_lineno  
        line = linecache.getline(filename, lineno)  
  
        print("%s(%d): %s" % (bname, lineno, line),  
              end=' ')  
  
    return self.localtrace
```





# **How do I register tracer?**

# Use `sys.settrace()`

```
import sys
from trace import my_tracing_func

def child():
    print('child')

def parent():
    child()

if __name__ == '__main__':
    sys.settrace(my_tracing_func)
    parent()
```



# Why Return a Tracer Function?



# Why Return a Tracer Function?

- ▶ You can have multiple tracer functions!
  - ▶ "System" Tracer
  - ▶ "Local" Tracer



# When does each tracer run?

1. System tracer runs for "Call" event.
2. Local tracer runs for all other events.

# Frame Attributes (again!)

```
>>> help(sys._getframe())
```

```
class frame(object)
```

```
...
```

```
| f_back          # this frame's caller  
| f_builtins      # builtins seen by this frame  
| f_code          # code object being executed  
| f_globals       # globals seen by this frame  
| f_lasti         # index of last instruction  
| f_lineno        # current line number in source  
| f_locals        # locals seen by this frame  
| f_trace        # tracing function for this frame
```



# **System and Local Example**

# System and Local Tracer

```
def system_tracer(frame, event, arg):  
    if event == 'call':  
        func_name = frame.f_code.co_name  
        print('... system_tracer: %r' % func_name)  
    return local_tracer  
  
def local_tracer(frame, event, arg):  
    if event == 'line':  
        lineno = frame.f_lineno  
        print('local_tracer: lineno %d' % lineno)  
    return local_tracer
```



# System and Local Tracer

```
def system_tracer(frame, event, arg):  
    if event == 'call':  
        func_name = frame.f_code.co_name  
        print('... system_tracer: %r' % func_name)  
    return local_tracer
```

```
def local_tracer(frame, event, arg):  
    if event == 'line':  
        lineno = frame.f_lineno  
        print('local_tracer: lineno %d' % lineno)  
    return local_tracer
```

# Example: count.py

```
def count_one():  
    print('1')  
    return 'Done'  
  
def count_two():  
    print('one')  
    print('two')  
    return 'Done'  
  
if __name__ == '__main__':  
    sys.settrace(system_tracer)  
    count_one()  
    count_two()  
    sys.settrace(None)
```

# When does each tracer run?

## Code

```
def count_one():  
    print('1')  
    return 'Done'
```

```
def count_two():  
    print('one')  
    print('two')  
    return 'Done'
```

## Event

Call

## Tracer

System

# When does each tracer run?

<u>Code</u>	<u>Event</u>	<u>Tracer</u>
<pre>def count_one():     <u>print('1')</u>     return 'Done'</pre>	Line	Local
<pre>def count_two():     print('one')     print('two')     return 'Done'</pre>		

# When does each tracer run?

## Code

```
def count_one():  
    print('1')  
    return 'Done'
```

```
def count_two():  
    print('one')  
    print('two')  
    return 'Done'
```

## Event

Return

## Tracer

Local

# When does each tracer run?

## Code

```
def count_one():  
    print('1')  
    return 'Done'
```

```
def count_two():  
    print('one')  
    print('two')  
    return 'Done'
```

## Event

Call

## Tracer

System

# When does each tracer run?

Code

Event

Tracer

```
def count_one():  
    print('1')  
    return 'Done'
```

```
def count_two():  
    print('one')  
    print('two')  
    return 'Done'
```

Line

Local

# When does each tracer run?

Code

Event

Tracer

```
def count_one():  
    print('1')  
    return 'Done'
```

```
def count_two():  
    print('one')  
    print('two')  
    return 'Done'
```

Line

Local



# When does each tracer run?

Code

Event

Tracer

```
def count_one():  
    print('1')  
    return 'Done'
```

```
def count_two():  
    print('one')  
    print('two')  
    return 'Done'
```

Return

Local



# Local Tracer can be None

If tracer returns None from "Call" event,  
no further tracing in frame.



# Recap: How tracing works

## 1. Create System tracer

- ▶ Signature: (frame, event, arg)
- ▶ Return local tracer, or None

## 2. Register System tracer

- ▶ `sys.settrace(system_tracer)`



# Agenda

1. Simple Tracing
2. How Does Tracing Work?
- 3. Custom Tracing**
4. Who Else Uses Custom Tracing?
5. Full-Fledged Example



# **Example #1**

## **Print Full Filename Tracer**

# Print Full Filename Tracer

```
def filename_tracer(frame, event, arg):
    fname = frame.f_code.co_filename
    co_name = frame.f_code.co_name
    num = frame.f_lineno
    line = linecache.getline(fname, num)

    if event == 'call':
        logger.info("--- %s: %s", fname, co_name)
    elif event == 'line':
        logger.info("%s(%d): %s", fname, num, line)

    return filename_tracer
```

# Print Full Filename Tracer

```
def filename_tracer(frame, event, arg):  
    fname = frame.f_code.co_filename  
    co_name = frame.f_code.co_name  
    num = frame.f_lineno  
    line = linecache.getline(fname, num)  
  
    if event == 'call':  
        logger.info("--- %s: %s", fname, co_name)  
    elif event == 'line':  
        logger.info("%s(%d): %s", fname, num, line)  
  
    return filename_tracer
```

# Print Full Filename Tracer

```
def filename_tracer(frame, event, arg):  
    fname = frame.f_code.co_filename  
    co_name = frame.f_code.co_name  
    num = frame.f_lineno  
    line = linecache.getline(fname, num)  
  
    if event == 'call':  
        logger.info("--- %s: %s", fname, co_name)  
    elif event == 'line':  
        logger.info("%s(%d): %s", fname, num, line)  
  
    return filename_tracer
```



# teapot\_filename\_trace.py

```
import sys
import requests
from filename_trace import filename_tracer

def teapot():
    url = 'http://httpbin.org/status/418'
    resp = requests.get(url)
    print(resp.status_code, resp.reason)
    print(resp.text)


if __name__ == '__main__':
    sys.settrace(filename_tracer)
    teapot()
    sys.settrace(None)
```



# Standard Library

```
$ grep client.py logs/trace.log | head -5
```

```
--- ../../lib/python3.5/http/client.py: __init__  
../../lib/python3.5/http/client.py(728):      self.timeout...  
../../lib/python3.5/http/client.py(729):      self.source_...  
../../lib/python3.5/http/client.py(730):      self.sock...  
../../lib/python3.5/http/client.py(731):      self._buffer...
```



# Requests

```
$ grep api.py logs/trace.log | head -5
```

```
--- ../../lib/python3.5/site-packages/requests/api.py: get  
../../lib/python3.5/site-packages/requests/api.py(69):      kwargs...  
../../lib/python3.5/site-packages/requests/api.py(70):      return...  
--- ../../lib/python3.5/site-packages/requests/api.py: request  
../../lib/python3.5/site-packages/requests/api.py(55):      with...
```



# **Example #2**

## **Print Function Arguments**

# Print function args

```
def trace_call_args(frame, event, arg):  
    if event == 'call':  
        code = frame.f_code  
        arg_count = code.co_argcount  
        arg_names = code.co_varnames[:arg_count]  
  
        print("%s: %s" % (code.co_filename, code.co_name))  
  
        for name in arg_names:  
            val = frame.f_locals[name]  
            print('{:<4}{}: {}'.format(' ', name, val))  
  
    return None
```

# Print function args

```
def trace_call_args(frame, event, arg):  
    if event == 'call':  
        code = frame.f_code  
        arg_count = code.co_argcount  
        arg_names = code.co_varnames[:arg_count]  
  
        print("%s: %s" % (code.co_filename, code.co_name))  
  
        for name in arg_names:  
            val = frame.f_locals[name]  
            print('{:<4}{}: {}'.format(' ', name, val))  
  
    return None
```

# Print function args

```
def trace_call_args(frame, event, arg):
    if event == 'call':
        code = frame.f_code
        arg_count = code.co_argcount
        arg_names = code.co_varnames[:arg_count]

        print("%s: %s" % (code.co_filename, code.co_name))

        for name in arg_names:
            val = frame.f_locals[name]
            print('{:<4}{}: {}'.format(' ', name, val))

    return None
```

# Example: add.py

```
import sys
from trace_call_args import trace_call_args

def add(x, y):
    return x + y

def add_with_defaults(x=49, y=50):
    return x + y

if __name__ == '__main__':
    sys.settrace(trace_call_args)
    add(1, 2)
    add_with_defaults()
    sys.settrace(None)
```





# Output: add.py

```
$ python add.py
```

```
add.py: add
```

```
    x: 1
```

```
    y: 2
```

```
add.py: add_with_defaults
```

```
    x: 49
```

```
    y: 50
```

# Example: var\_args\_1.py

```
import sys
from trace_call_args import trace_call_args

def var_args(*args, **kwargs):
    x = 1
    y = 2
    return x, y

if __name__ == '__main__':
    sys.settrace(trace_call_args)
    args = (1, 2, 'foo')
    kwargs = {'a': 8, 'b': 9}
    var_args(*args, **kwargs)
    sys.settrace(None)
```



# Output: `var_args_1.py`

```
$ python var_args_1.py
```

```
var_args_1.py: var_args
```

# Print function args

```
def trace_call_args(frame, event, arg):
    if event == 'call':
        code = frame.f_code
        arg_count = code.co_argcount
        arg_names = code.co_varnames[:arg_count]

        print("%s: %s" % (code.co_filename, code.co_name))

        for name in arg_names:
            val = frame.f_locals[name]
            print('{:<4}{}: {}'.format(' ', name, val))

    return None
```

# Print function args: Take 2

```
def trace_call_args_2(frame, event, arg):  
    if event == 'call':  
        code = frame.f_code  
  
        print("%s: %s" % (code.co_filename, code.co_name))  
  
        for name in frame.f_locals:  
            val = frame.f_locals[name]  
            print('{:<4}{:<7}: {}'.format(' ', name, val))  
  
    return None
```

# Example: var\_args\_2.py

```
import sys
from trace_call_args import trace_call_args_2

def var_args(*args, **kwargs):
    x = 1
    y = 2
    return x, y

if __name__ == '__main__':
    sys.settrace(trace_call_args_2)
    args = (1, 2, 'foo')
    kwargs = {'a': 8, 'b': 9}
    var_args(*args, **kwargs)
    sys.settrace(None)
```



# Output: var\_args\_2.py

```
$ python var_args_2.py
```

```
var_args_2.py: var_args  
    args      : (1, 2, 'foo')  
    kwargs    : {'b': 9, 'a': 8}
```



# Agenda

1. Simple Tracing
2. How Does Tracing Work?
3. Custom Tracing
- 4. Who Else Uses Custom Tracing?**
5. Full-Fledged Example





# coverage.py

- ▶ Pytracer
- ▶ CTracer

# coverage.py tracer

```
class PyTracer(object):
    ...
    def _trace(self, frame, event, arg_unused):
        """
        The trace function passed to sys.settrace.
        """
        ...
        if event == 'call':
            ...
        elif event == 'line':
            ...
        elif event == 'return':
            ...
        return self._trace
```



# Python Debugger (Pdb)

- ▶ Built on top of Bdb
- ▶ "Debugger Framework"



# Bdb tracer

```
class Bdb:
```

```
    ...
```

```
    def trace_dispatch(self, frame, event, arg):
```

```
        ...
```

```
        if event == 'line':
```

```
            return self.dispatch_line(frame)
```

```
        if event == 'call':
```

```
            return self.dispatch_call(frame, arg)
```

```
        if event == 'return':
```

```
            return self.dispatch_return(frame, arg)
```

```
        if event == 'exception':
```

```
            return self.dispatch_exception(frame, arg)
```

```
        ...
```



# pythontutor.com

"PGLogger is a subclass of **bdb.Bdb**...

Here's where the magic happens...

As the user's program is running, **bdb** will pause execution at every function call, return, exception, and single-line step (most common)."



# Agenda

1. Simple Tracing
2. How Does Tracing Work?
3. Custom Tracing
4. Who Else Uses Custom Tracing?
- 5. Full-Fledged Example**



**Let's build a Trace Visualizer...**

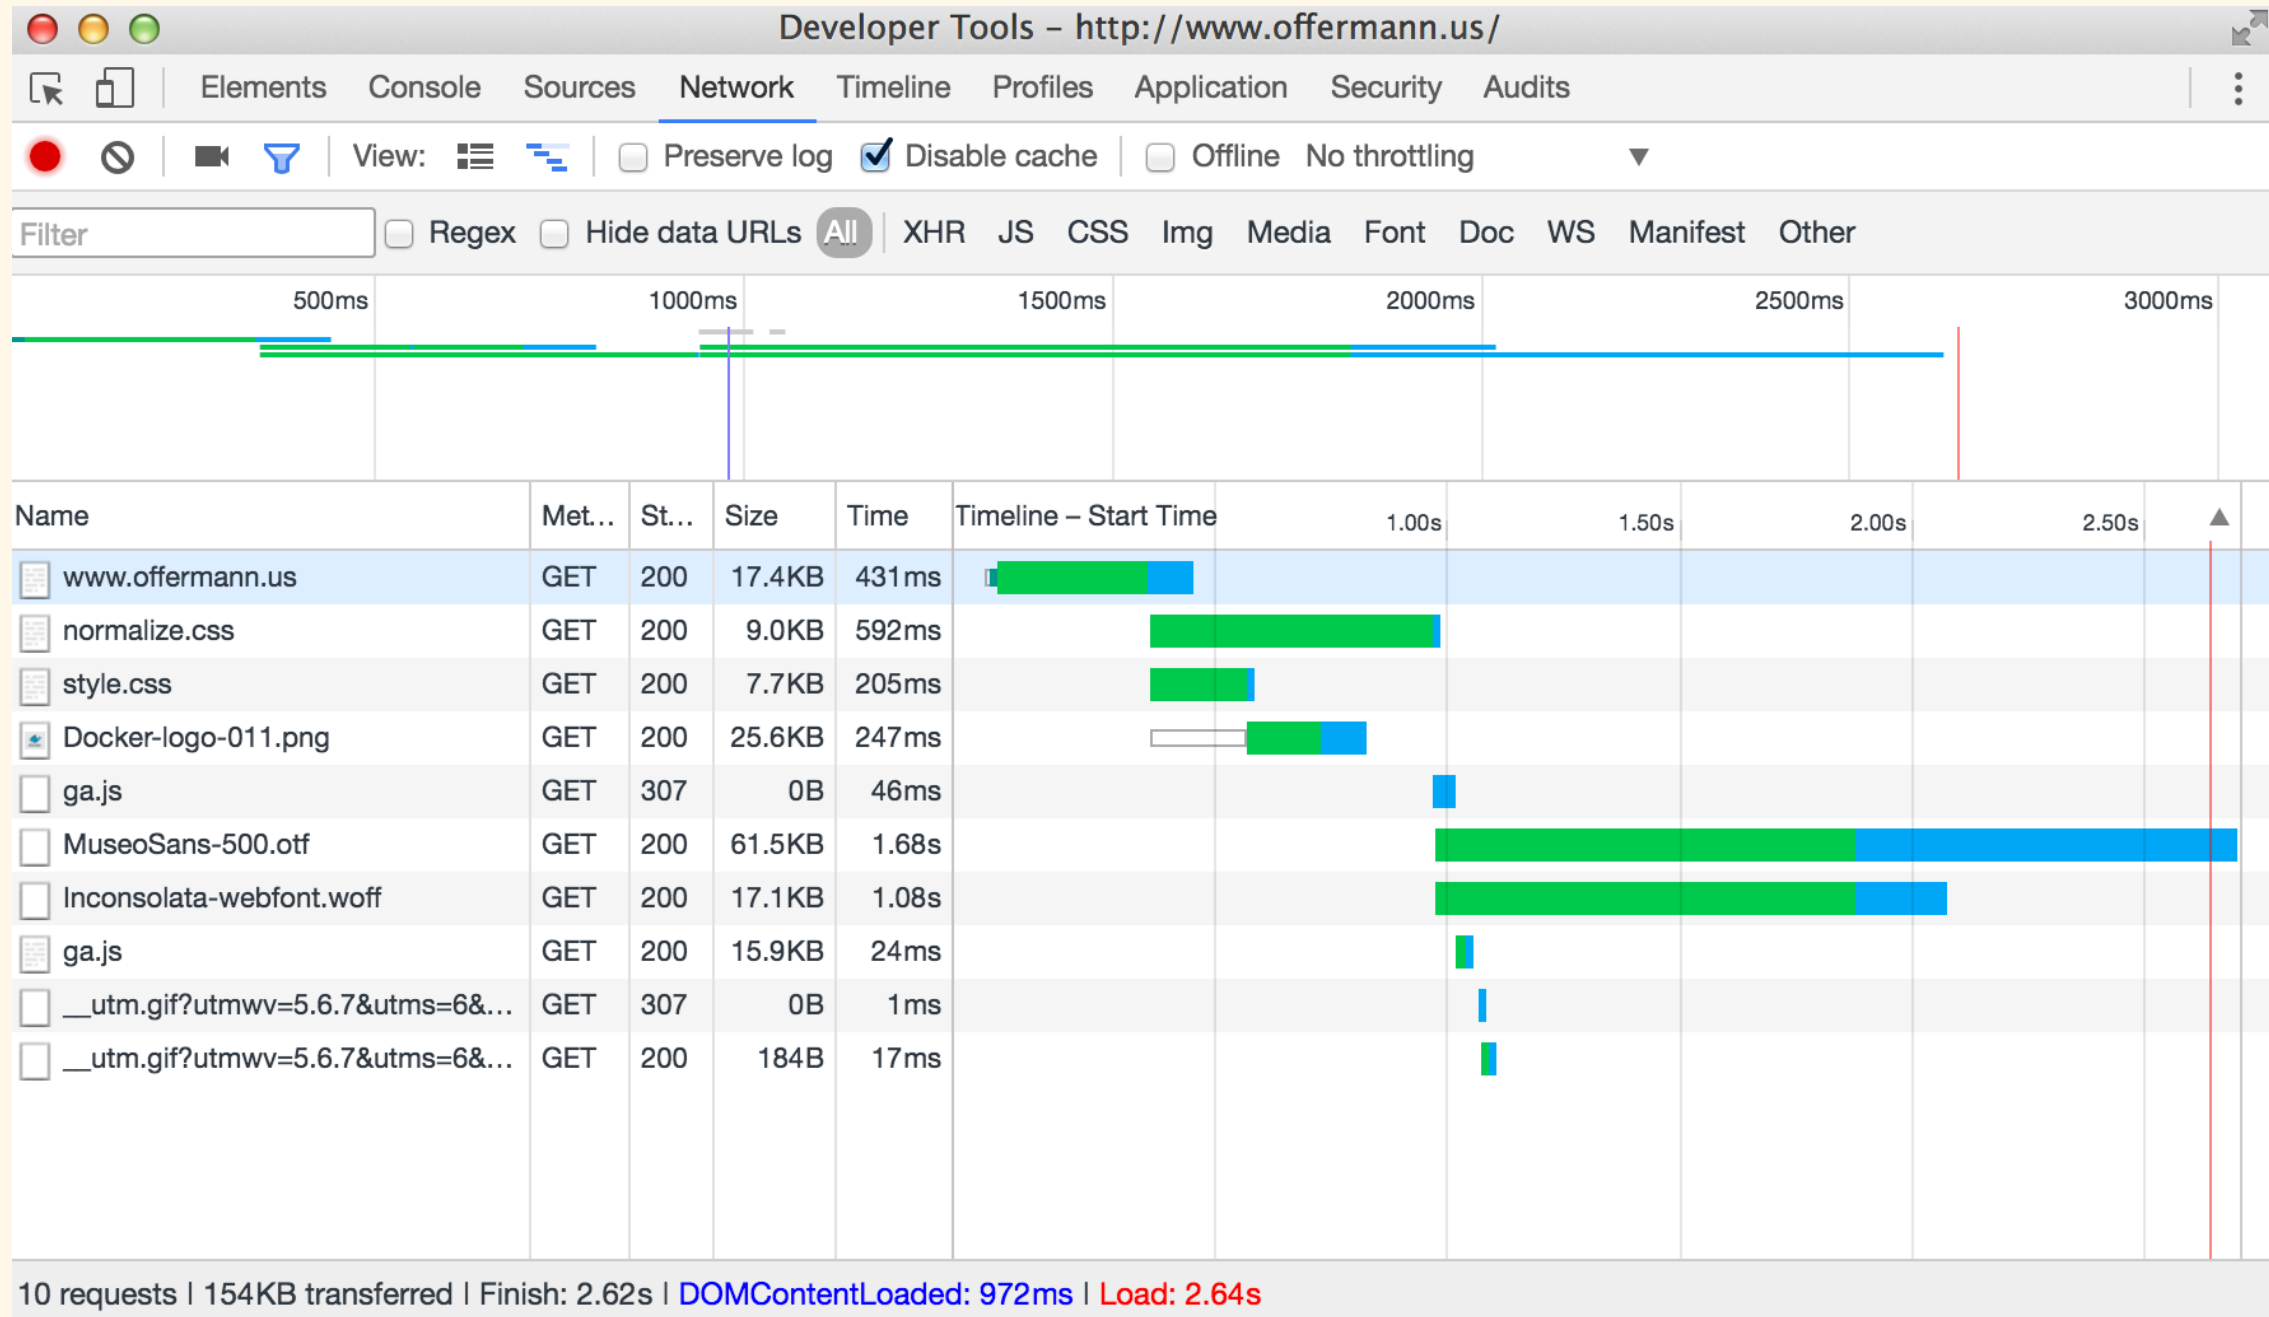


**Let's build a Trace Visualizer...**

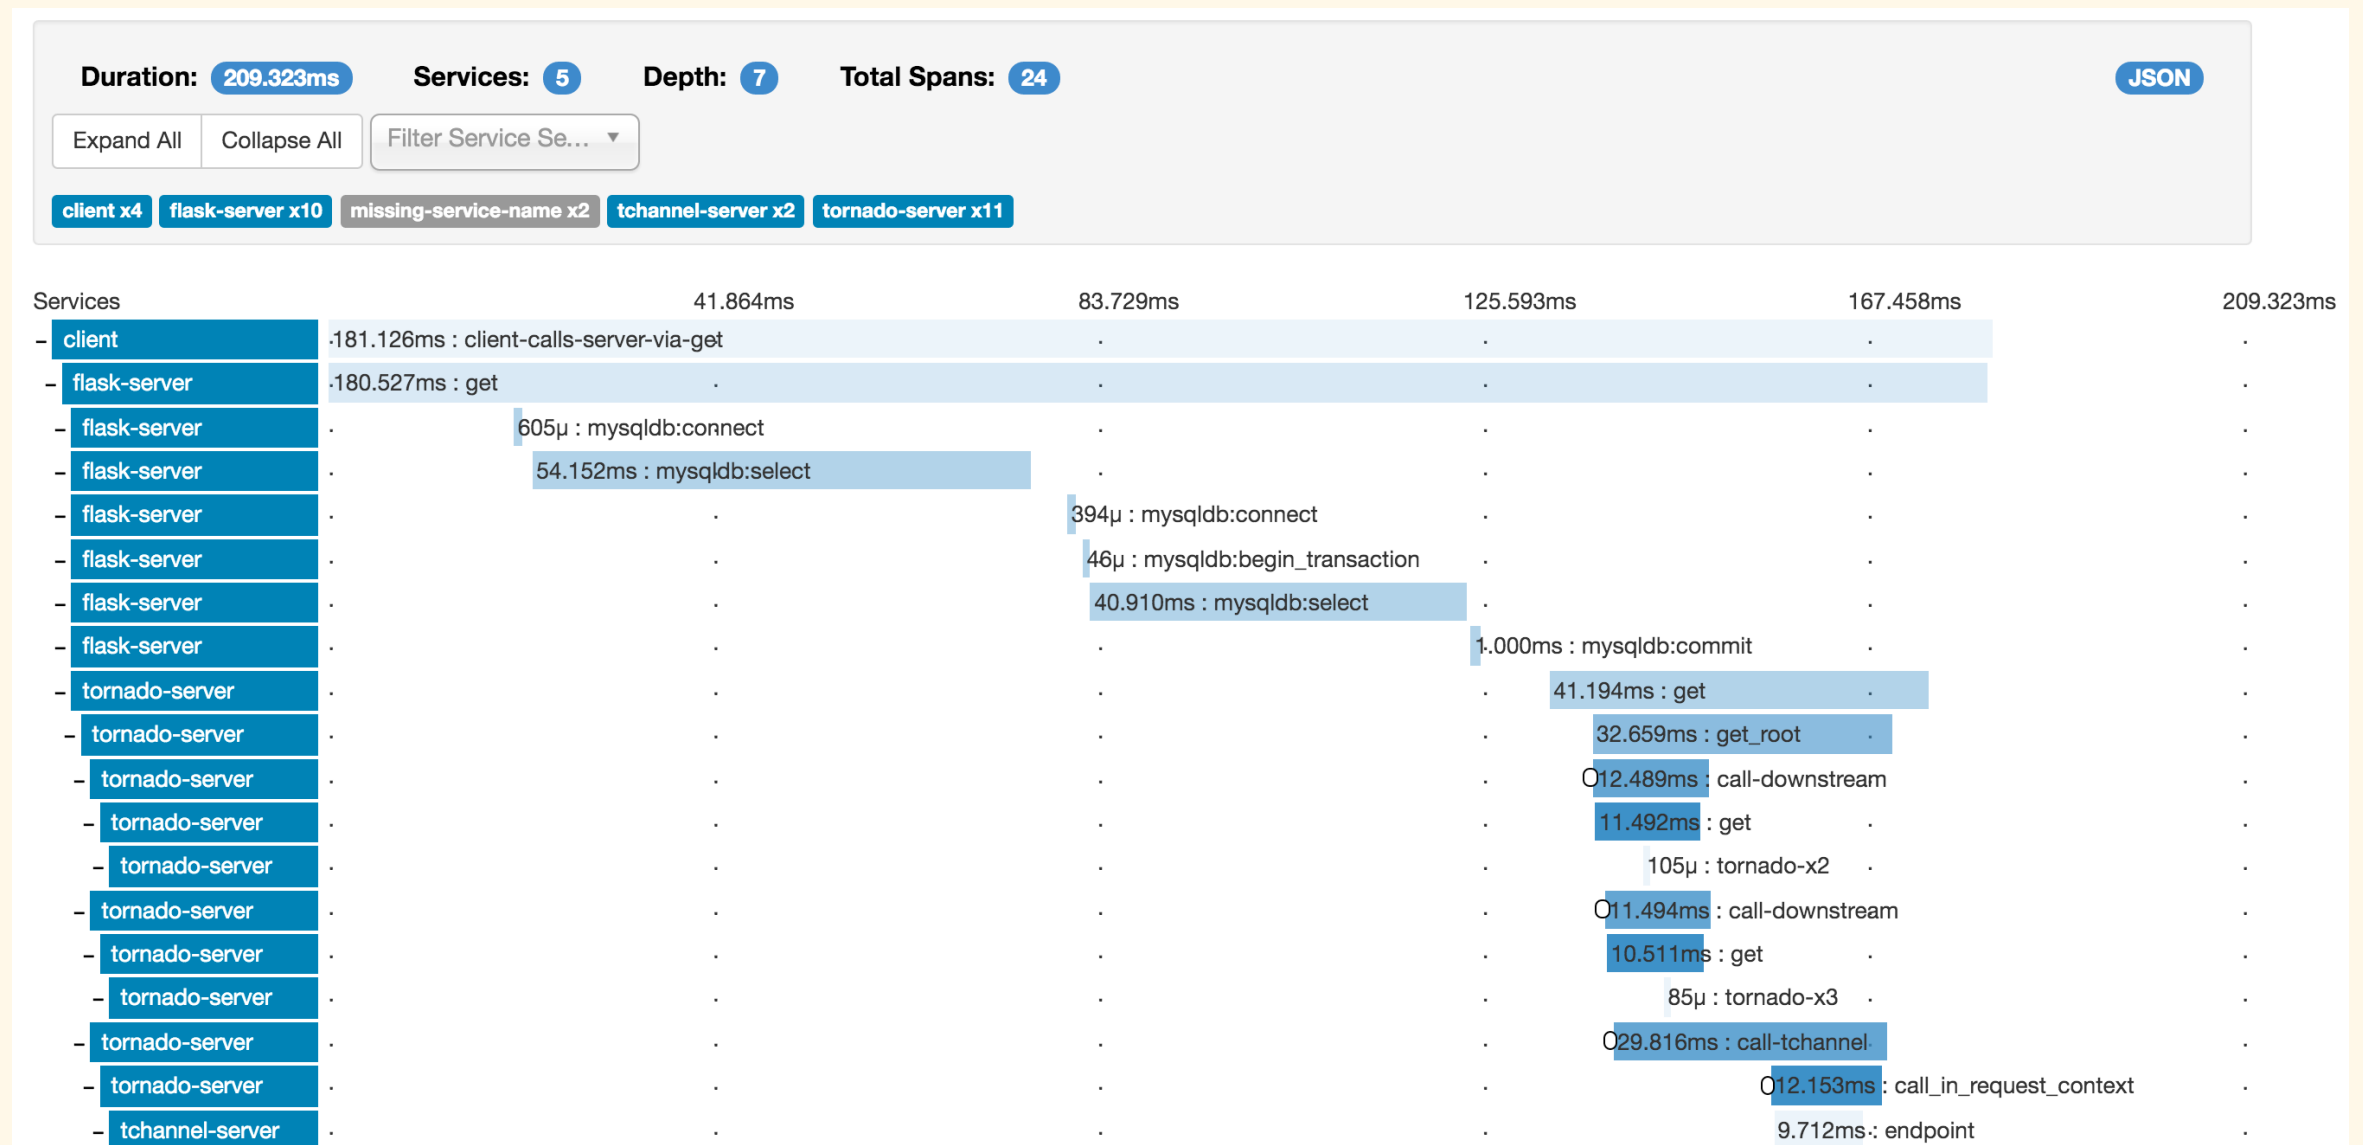
**Or, borrow one!**



# Chrome Dev Tools



# Zipkin





# What is Zipkin?

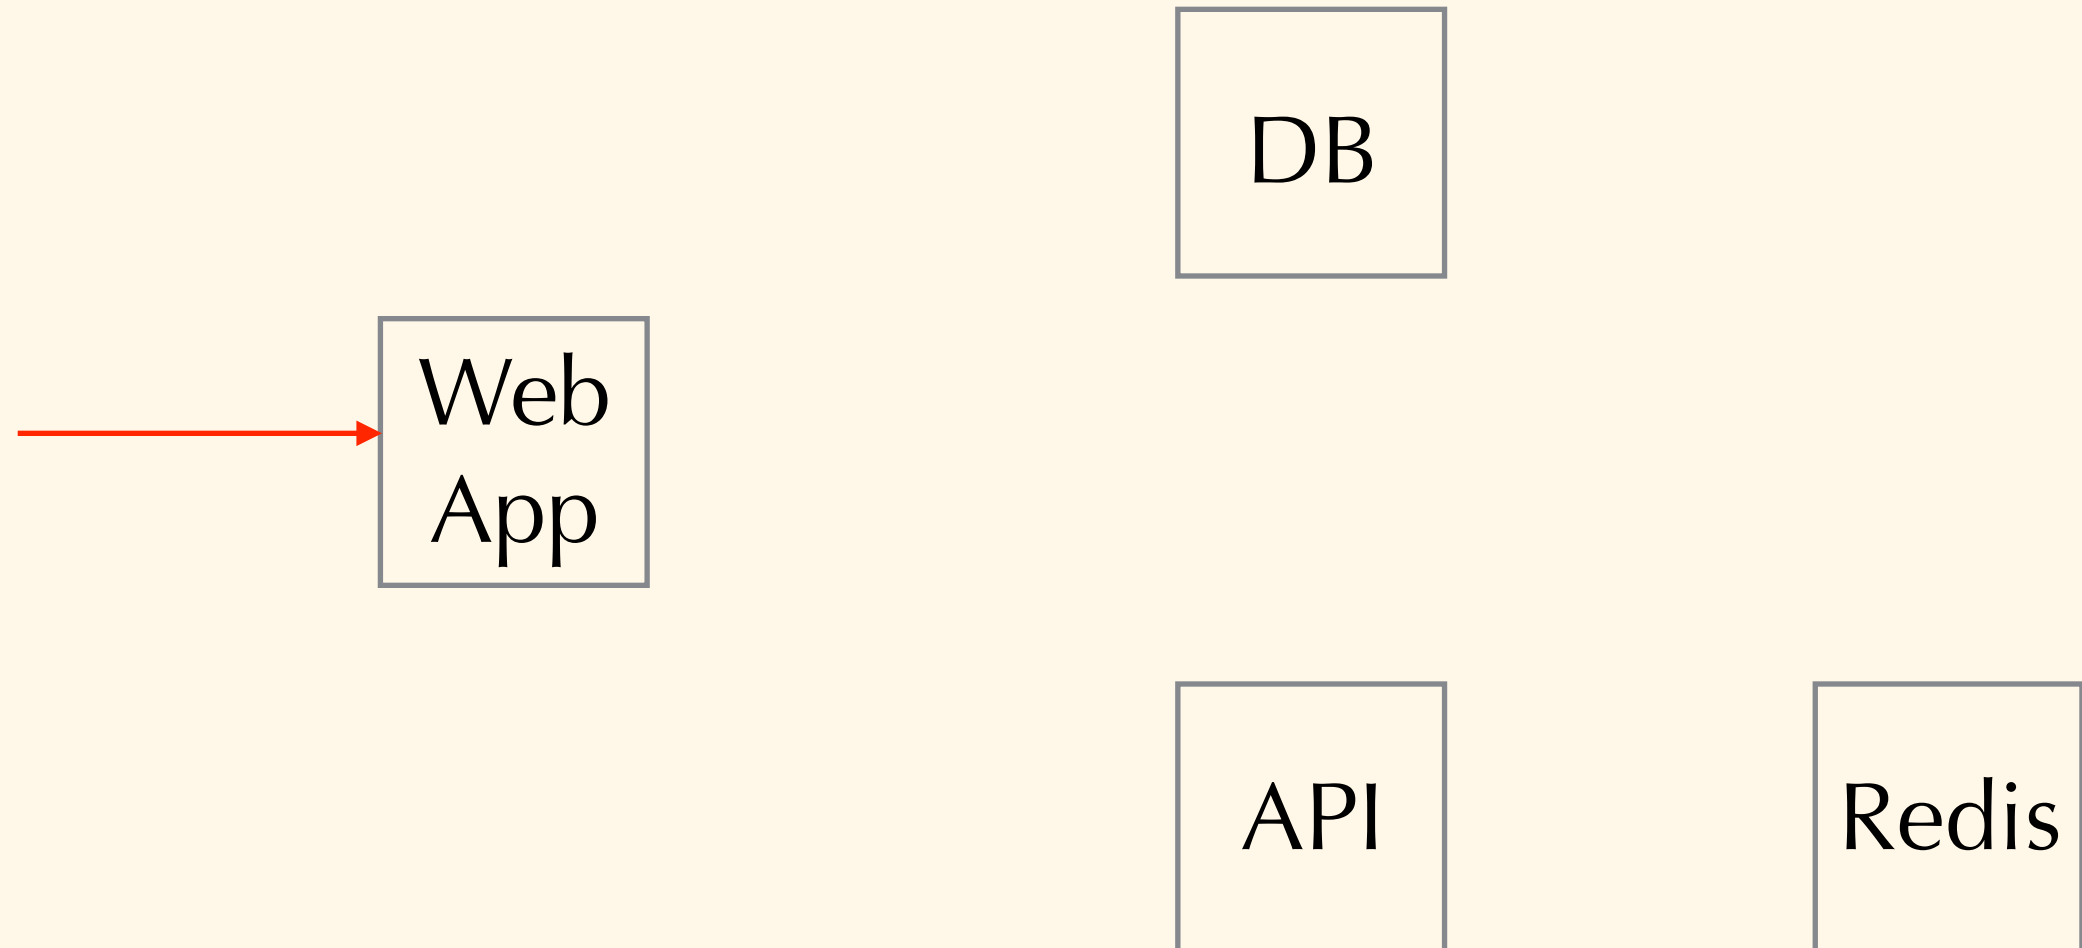
- ▶ Distributed Tracing for microservices
- ▶ Many RPC calls to serve single request.
  - ▶ HTTP calls to API endpoints
  - ▶ Database queries



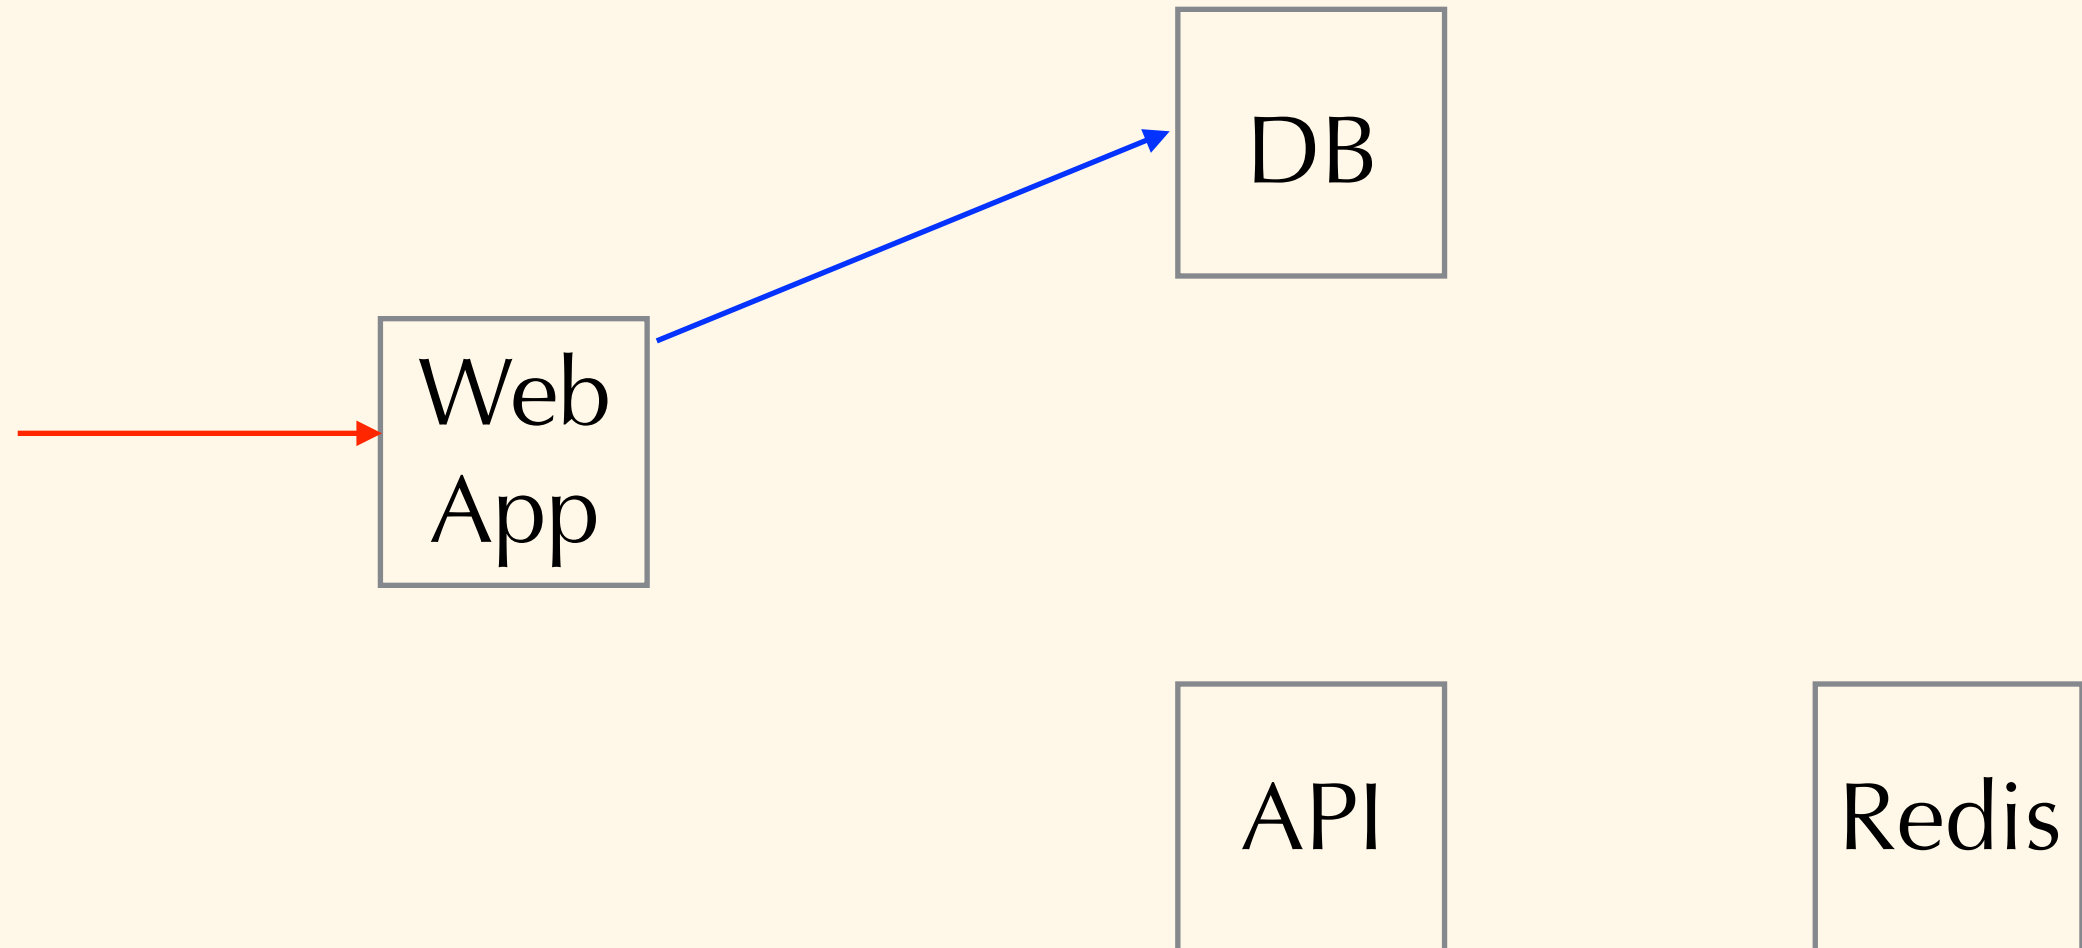
# What is Zipkin? (2)

- ▶ **Spans:** Individual RPC calls.
- ▶ **Trace:** Collection of Spans.


# Zipkin Example




# Zipkin Example



# Zipkin Example




# Zipkin Example







# Zipkin Example




# Zipkin Example




# Zipkin Example



# Zipkin Example



# Zipkin Diagram






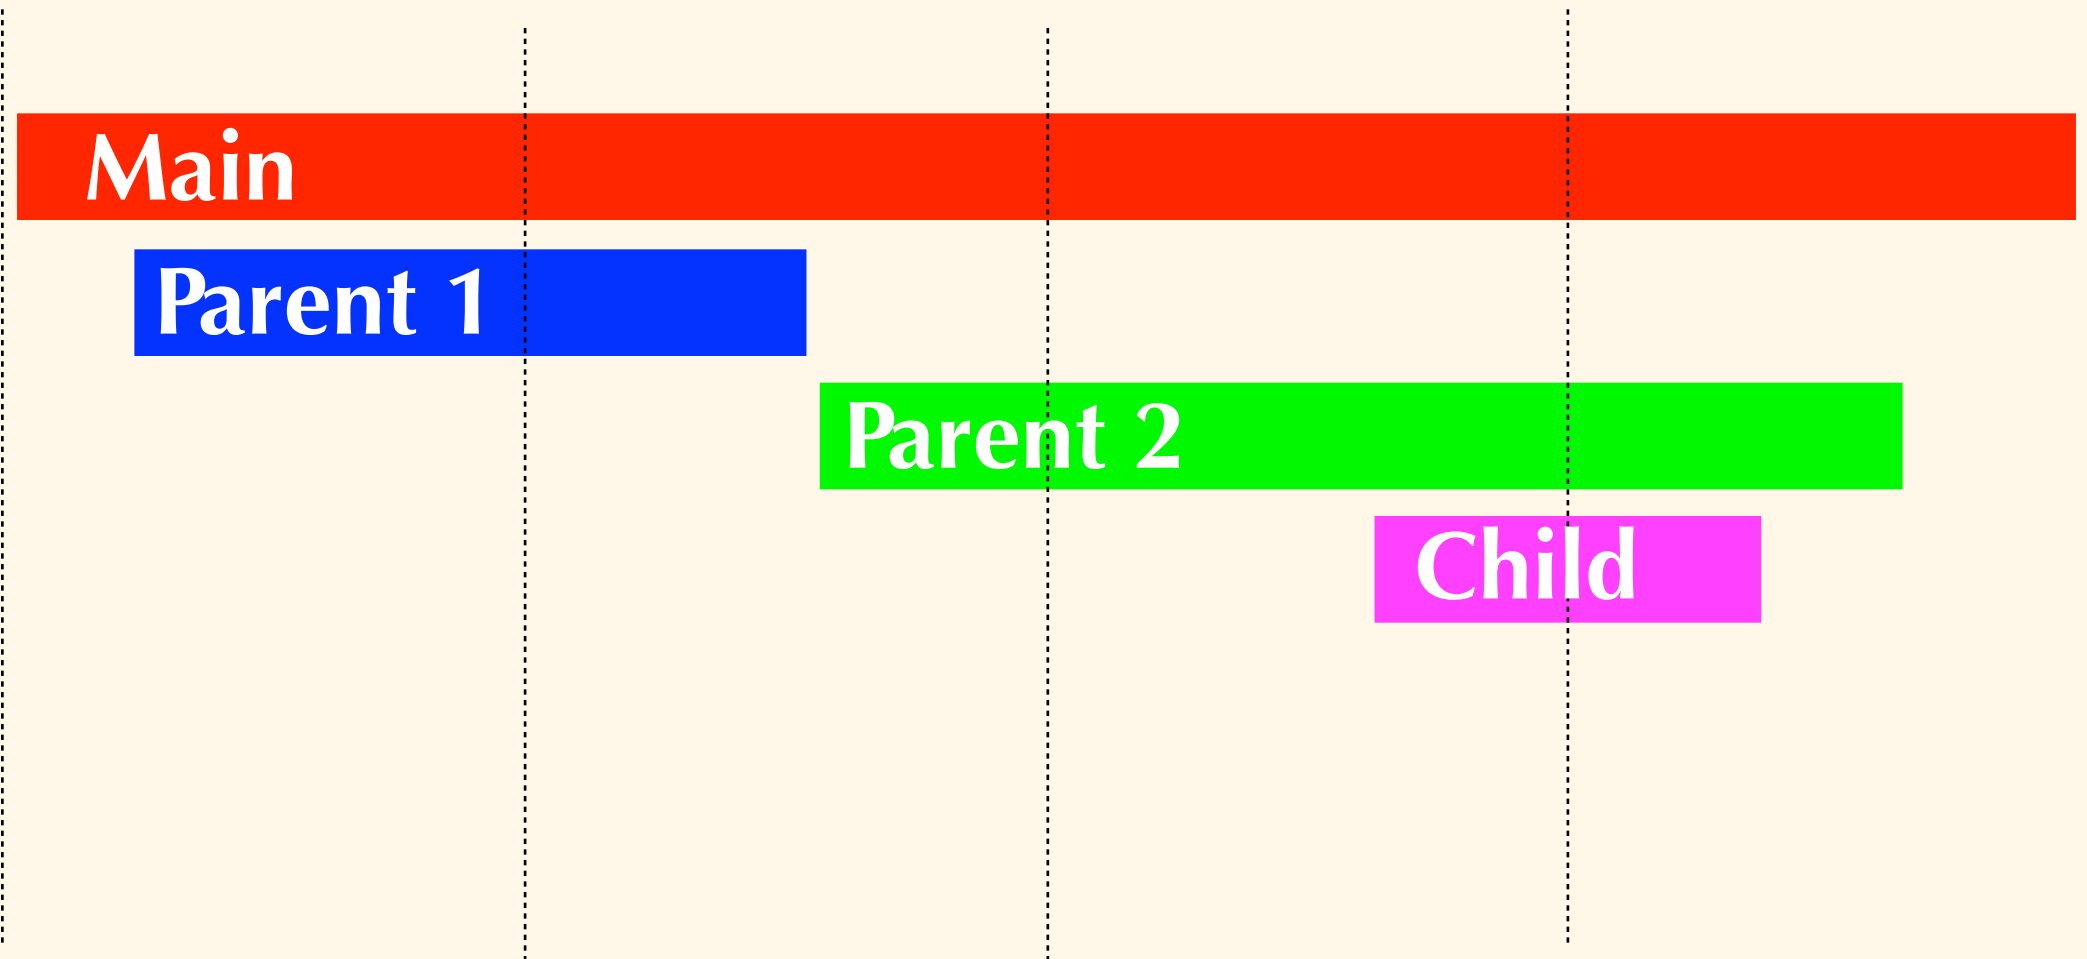
# Translation to Tracing

- ▶ Trace -> Run Program
- ▶ Span -> Function call (frame)

# Translation to Tracing



# Spans = Frames





# Zipkin Span Format

```
[{
  "parentId": "0ed0c2af926048a7",
  "name": "child",
  "timestamp": 1475272253175560.0,
  "debug": true,
  "id": "0e8a6ecd28d0b1bd",
  "annotations": [
    {
      "endpoint": {
        "ipv4": "127.0.0.1",
        "port": 8888,
        "serviceName": "simple_service"
      },
      "timestamp": 1475272253175560.0,
      "value": "sr"
    }
  ],
  "traceId": "22bf9db1413e9fcf"
}]
```

# Zipkin Span Format

```
[{
  "parentId": "0ed0c2af926048a7",
  "name": "child",
  "timestamp": 1475272253175560.0,
  "debug": true,
  "id": "0e8a6ecd28d0b1bd",
  "annotations": [
    {
      "endpoint": {
        "ipv4": "127.0.0.1",
        "port": 8888,
        "serviceName": "simple_service"
      },
      "timestamp": 1475272253175560.0,
      "value": "sr"
    }
  ],
  "traceId": "22bf9db1413e9fcf"
}]
```

# Zipkin Span Format

```
[{
  "parentId": "0ed0c2af926048a7",
  "name": "child",
  "timestamp": 1475272253175560.0,
  "debug": true,
  "id": "0e8a6ecd28d0b1bd",
  "annotations": [
    {
      "endpoint": {
        "ipv4": "127.0.0.1",
        "port": 8888,
        "serviceName": "simple_service"
      },
      "timestamp": 1475272253175560.0,
      "value": "sr"
    }
  ],
  "traceId": "22bf9db1413e9fcf"
}]
```

# Zipkin Span Format

```
[{
  "parentId": "0ed0c2af926048a7",
  "name": "child",
  "timestamp": 1475272253175560.0,
  "debug": true,
  "id": "0e8a6ecd28d0b1bd",
  "annotations": [
    {
      "endpoint": {
        "ipv4": "127.0.0.1",
        "port": 8888,
        "serviceName": "simple_service"
      },
      "timestamp": 1475272253175560.0,
      "value": "sr"
    }
  ],
  "traceId": "22bf9db1413e9fcf"
}]
```

# Zipkin Span Format

```
[{
  "parentId": "0ed0c2af926048a7",
  "name": "child",
  "timestamp": 1475272253175560.0,
  "debug": true,
  "id": "0e8a6ecd28d0b1bd",
  "annotations": [
    {
      "endpoint": {
        "ipv4": "127.0.0.1",
        "port": 8888,
        "serviceName": "simple_service"
      },
      "timestamp": 1475272253175560.0,
      "value": "sr"
    }
  ],
  "traceId": "22bf9db1413e9fcf"
}]
```



# Plan

1. Trace program
2. Log results
3. Convert log to Zipkin JSON
4. Visualize!

# Tracer

```
class ZipkinTracer(object):

    def __init__(self):
        self.span_ids = {} # {frame: span_id}

    def __call__(self, frame, event, arg):
        if event == 'call':
            return self.trace_call(frame, event, arg)
        elif event == 'return':
            return self.trace_return(frame, event, arg)
        else:
            return self
```

# Tracer for Call Event

```
def trace_call(self, frame, event, arg):  
    current_id = os.urandom(8).hex()  
    self.span_ids[frame] = current_id  
  
    parent_id = self.parent_span_id(frame)  
    func_name = frame.f_code.co_name  
    zipkin_logger.debug('%s,%s,%s,%f,%s' % (  
        current_id, parent_id, func_name,  
        time.time(), 'sr'))  
  
    return self
```



# Tracer for Return Event

```
def trace_return(self, frame, event, arg):
    current_id = self.span_ids[frame]
    parent_id = self.parent_span_id(frame)
    func_name = frame.f_code.co_name
    zipkin_logger.debug('%s,%s,%s,%f,%s' % (
        current_id, parent_id, func_name,
        time.time(), 'ss'))

    del self.span_ids[frame]
    return self
```



# Tracer

```
def parent_span_id(self, frame):  
    parent = frame.f_back  
    return self.span_ids.get(parent, '')
```



# Log File

```
3f5edd6b77ed0bf5,,main,1475053714.172785,sr
b6c53ba3caa2f504,3f5edd6b77ed0bf5,parent,1475053714.172898,sr
61b79578737e7423,b6c53ba3caa2f504,child,1475053714.172950,sr
61b79578737e7423,b6c53ba3caa2f504,child,1475053714.173023,ss
09abc624393d70d5,b6c53ba3caa2f504,child,1475053714.173066,sr
09abc624393d70d5,b6c53ba3caa2f504,child,1475053714.173134,ss
5d9be223d8d7bd4c,b6c53ba3caa2f504,child,1475053714.173184,sr
5d9be223d8d7bd4c,b6c53ba3caa2f504,child,1475053714.173259,ss
b6c53ba3caa2f504,3f5edd6b77ed0bf5,parent,1475053714.173304,ss
3f5edd6b77ed0bf5,,main,1475053714.173345,ss
```



# log2span.py

```
for line in read_log(filename):  
    span = make_span(*line)  
    send_span(span)
```

# make\_span

```
def make_span(span_id, parent_id, func_name, timestamp,
               annotation_value, endpoint=DEFAULT_ENDPOINT, debug=True):

    span = Span(
        span_id=span_id,
        parent_id=parent_id,
        trace_id=DEFAULT_TRACE_ID,
        name=func_name,
        timestamp= timestamp * (10 ** 6),
        debug=debug
    )

    annotation = Annotation(
        endpoint=DEFAULT_ENDPOINT,
        timestamp= timestamp * (10 ** 6),
        value = annotation_value
    )

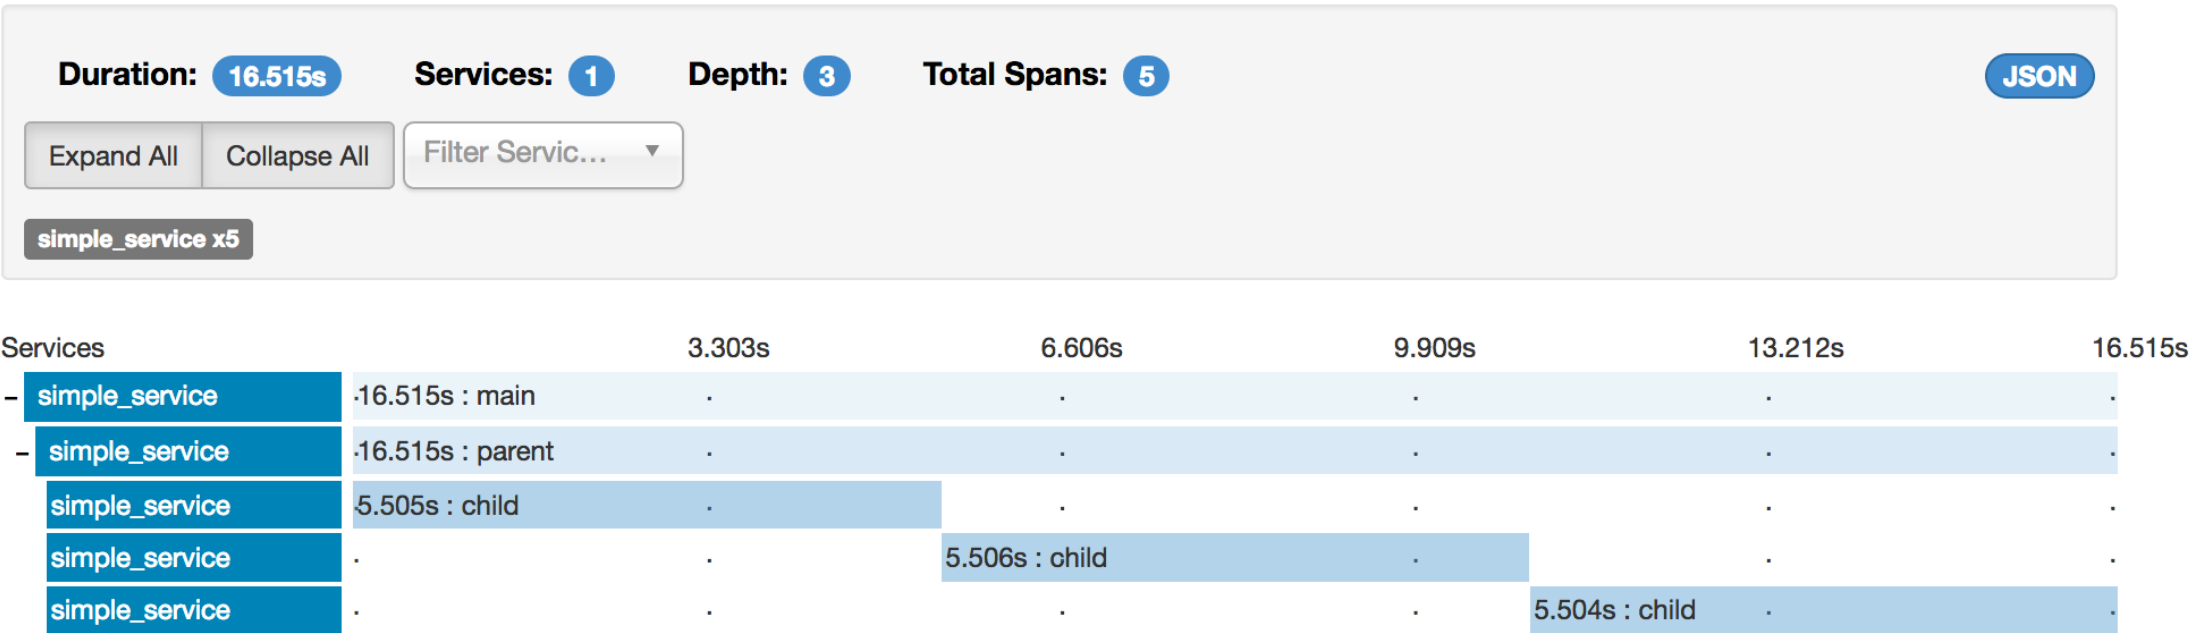
    span.add_annotation(annotation)
    return span
```



# send\_span

```
def send_span(span):  
    url = 'http://192.168.99.100:9411/api/v1/spans'  
    payload = [span.asdict()]  
    resp = requests.post(url, json=payload)  
    print(resp)
```

# End Result





# Conclusion

- ▶ Nothing scary about trace functions.
- ▶ Tracing can be extraordinarily verbose.
- ▶ Capable of deep introspection.






# Conclusion

- ▶ Nothing scary about trace functions.
- ▶ Tracing can be extraordinarily verbose.
- ▶ Capable of deep introspection.
- ▶ **Repurposing UI is OK!**



# Further Reading

- ▶ [\*\*github.com/toffer/talk-custom-tracing\*\*](https://github.com/toffer/talk-custom-tracing)
  - ▶ Code Examples
  - ▶ Slide Deck



**Thanks!**